

Automating Data Exploration with R

Basic Exploration

Pair-wise correlations

Now that we have all our data transformed into numbers, we're ready to take a deeper look at it. Here we'll build two functions, a pair-wise correlation function and a more complicated function to get the p-value and correlation for each feature pair in the data.

But before we jump in, let's take a quick look at correlations. A great way to explore new data is to use a pairwise correlation matrix. This will measure the correlation between every combination of your variables. It doesn't really matter if you have an outcome (or response) variable at this point, it will compare everything against everything else.

For those not familiar with the correlation coefficient, it is simply a measure of similarity between two vectors of numbers. The measure value can range between 1 and -1, where 1 is perfectly correlated, -1 is perfectly inversely correlated, and 0 is not correlated at all:

```
print(cor(1:5,1:5))
```

```
## [1] 1
```

```
print(cor(1:5,seq(100,500,100)))
```

```
## [1] 1
```

```
print(cor(1:5,5:1))
```

```
## [1] -1
```

```
print(cor(1:5,c(1,2,3,4,4)))
```

```
## [1] 0.9701425
```

We'll use the `mtcars` data set that is already included in the **R** base package - it has the advantage of being fully numeric and clean.

```
# install.packages('dplyr')  
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'  
##  
## The following objects are masked from 'package:stats':  
##  
##     filter, lag  
##  
## The following objects are masked from 'package:base':  
##  
##     intersect, setdiff, setequal, union
```

```
# install.packages('reshape2')  
library(reshape2)  
  
data_set <- mtcars  
d_cor <- as.matrix(cor(data_set))  
d_cor
```

```
##           mpg           cyl           disp           hp           drat           wt
## mpg      1.0000000 -0.8521620 -0.8475514 -0.7761684  0.68117191 -0.8676594
## cyl     -0.8521620  1.0000000  0.9020329  0.8324475 -0.69993811  0.7824958
## disp    -0.8475514  0.9020329  1.0000000  0.7909486 -0.71021393  0.8879799
## hp      -0.7761684  0.8324475  0.7909486  1.0000000 -0.44875912  0.6587479
## drat     0.6811719 -0.6999381 -0.7102139 -0.4487591  1.00000000 -0.7124406
## wt      -0.8676594  0.7824958  0.8879799  0.6587479 -0.71244065  1.0000000
## qsec     0.4186840 -0.5912421 -0.4336979 -0.7082234  0.09120476 -0.1747159
## vs       0.6640389 -0.8108118 -0.7104159 -0.7230967  0.44027846 -0.5549157
## am       0.5998324 -0.5226070 -0.5912270 -0.2432043  0.71271113 -0.6924953
## gear     0.4802848 -0.4926866 -0.5555692 -0.1257043  0.69961013 -0.5832870
## carb    -0.5509251  0.5269883  0.3949769  0.7498125 -0.09078980  0.4276059
##           qsec           vs           am           gear           carb
## mpg      0.41868403  0.6640389  0.59983243  0.4802848 -0.55092507
## cyl     -0.59124207 -0.8108118 -0.52260705 -0.4926866  0.52698829
## disp    -0.43369788 -0.7104159 -0.59122704 -0.5555692  0.39497686
## hp      -0.70822339 -0.7230967 -0.24320426 -0.1257043  0.74981247
## drat     0.09120476  0.4402785  0.71271113  0.6996101 -0.09078980
## wt      -0.17471588 -0.5549157 -0.69249526 -0.5832870  0.42760594
## qsec     1.00000000  0.7445354 -0.22986086 -0.2126822 -0.65624923
## vs       0.74453544  1.0000000  0.16834512  0.2060233 -0.56960714
## am      -0.22986086  0.1683451  1.00000000  0.7940588  0.05753435
## gear    -0.21268223  0.2060233  0.79405876  1.0000000  0.27407284
## carb    -0.65624923 -0.5696071  0.05753435  0.2740728  1.00000000
```

For example, if you compare the two left utmost columns, we see that `mpg` is negatively correlated to `cyl`. Remember, correlations range from 1 to -1, with 1 being absolutely positively correlated, -1 being absolutely negatively correlated and 0 showing no correlations at all.

| mpg | cyl |
|-----|----------------------|
| mpg | 1.0000000 -0.8521620 |

```
d_cor_melt <- arrange(melt(d_cor), -(value))

# clean up
pair_wise_correlation_matrix <- filter(d_cor_melt, Var1 != Var2)
pair_wise_correlation_matrix <- filter(pair_wise_correlation_matrix, is.na(value)==FALSE)


# remove pair dups
dim(pair_wise_correlation_matrix)
```

```
## [1] 110 3
```

```
pair_wise_correlation_matrix <- pair_wise_correlation_matrix[seq(1, nrow(pair_wise_correlation_matrix), by=2),]
dim(pair_wise_correlation_matrix)
```

```
## [1] 55 3
```

```
plot(pair_wise_correlation_matrix$value)
```



Let's build the above into a function:

```

Get_Fast_Correlations <- function(data_set, features_to_ignore=c(), size_cap=5000) {
  require(dplyr)
  require(reshape2)

  data_set <- data_set[,setdiff(names(data_set), features_to_ignore)]

  if (size_cap > nrow(data_set)) {
    data_set = data_set[sample(nrow(data_set), size_cap),]
  } else {
    data_set = data_set[sample(nrow(data_set), nrow(data_set)),]
  }

  d_cor <- as.matrix(cor(data_set))
  d_cor_melt <- arrange(melt(d_cor), -(value))

  # clean up
  pair_wise_correlation_matrix <- filter(d_cor_melt, Var1 != Var2)
  pair_wise_correlation_matrix <- filter(pair_wise_correlation_matrix, is.na(value)==FALSE)

  # remove pair dups
  dim(pair_wise_correlation_matrix)
  pair_wise_correlation_matrix <- pair_wise_correlation_matrix[seq(1, nrow(pair_wise_correlation_matrix), by=2),
]
  dim(pair_wise_correlation_matrix)

  plot(pair_wise_correlation_matrix$value)
  return(pair_wise_correlation_matrix)
}

```

We'll use the `psych` (<https://cran.r-project.org/web/packages/psych/index.html>) library to do a lot of the heavy lifting. We will delegate all the math to the `corr.test` function. It returns the following results (from the help files and plenty more there: `?corr.test`):

- **r**: The matrix of correlations
- **n**: Number of cases per correlation
- **t**: Value of t-test for each correlation
- **p**: Two tailed probability of t for each correlation
- **se**: Standard error of the correlation
- **ci**: The alpha/2 lower and upper values

```
# install.packages('psych')
library(psych)

data_set <- mtcars
featurenames_copy <- names(data_set)

# strip var names to index for pair wise identification
names(data_set) <- seq(1:ncol(data_set))
cor_data_df <- corr.test(data_set)

# apply var names to correlation matrix over index
rownames(cor_data_df$r) <- featurenames_copy
colnames(cor_data_df$r) <- featurenames_copy

names(cor_data_df)
```

```
## [1] "r"      "n"      "t"      "p"      "se"     "adjust" "sym"    "ci"
## [9] "Call"
```


```
# matrix of correlations
cor_data_df$r
```

```
##           mpg           cyl           disp           hp           drat           wt
## mpg      1.0000000 -0.8521620 -0.8475514 -0.7761684  0.68117191 -0.8676594
## cyl     -0.8521620  1.0000000  0.9020329  0.8324475 -0.69993811  0.7824958
## disp    -0.8475514  0.9020329  1.0000000  0.7909486 -0.71021393  0.8879799
## hp      -0.7761684  0.8324475  0.7909486  1.0000000 -0.44875912  0.6587479
## drat     0.6811719 -0.6999381 -0.7102139 -0.4487591  1.00000000 -0.7124406
## wt      -0.8676594  0.7824958  0.8879799  0.6587479 -0.71244065  1.0000000
## qsec     0.4186840 -0.5912421 -0.4336979 -0.7082234  0.09120476 -0.1747159
## vs       0.6640389 -0.8108118 -0.7104159 -0.7230967  0.44027846 -0.5549157
## am       0.5998324 -0.5226070 -0.5912270 -0.2432043  0.71271113 -0.6924953
## gear     0.4802848 -0.4926866 -0.5555692 -0.1257043  0.69961013 -0.5832870
## carb    -0.5509251  0.5269883  0.3949769  0.7498125 -0.09078980  0.4276059
##           qsec           vs           am           gear           carb
## mpg      0.41868403  0.6640389  0.59983243  0.4802848 -0.55092507
## cyl     -0.59124207 -0.8108118 -0.52260705 -0.4926866  0.52698829
## disp    -0.43369788 -0.7104159 -0.59122704 -0.5555692  0.39497686
## hp      -0.70822339 -0.7230967 -0.24320426 -0.1257043  0.74981247
## drat     0.09120476  0.4402785  0.71271113  0.6996101 -0.09078980
## wt      -0.17471588 -0.5549157 -0.69249526 -0.5832870  0.42760594
## qsec     1.00000000  0.7445354 -0.22986086 -0.2126822 -0.65624923
## vs       0.74453544  1.0000000  0.16834512  0.2060233 -0.56960714
## am      -0.22986086  0.1683451  1.00000000  0.7940588  0.05753435
## gear    -0.21268223  0.2060233  0.79405876  1.0000000  0.27407284
## carb    -0.65624923 -0.5696071  0.05753435  0.2740728  1.00000000
```

Let's visualize our correlation results using `cor.plot` from the `psych` (<https://cran.r-project.org/web/packages/psych/index.html>) library and `corrplot.mixed` from the `corrplot` (<https://cran.r-project.org/web/packages/corrplot/vignettes/corrplot-intro.html>) library.


```
cor.plot(cor_data_df$r)
```


Correlation plot



Or even prettier with the `corrplot` library:

```
#install.packages('corrplot')  
library(corrplot)  
corrplot.mixed(cor_data_df$r, lower="circle", upper="color",  
tl.pos="lt", diag="n", order="hclust", hclust.method="complete")
```



We're not going to add these plotting functions into our pipeline, but we will create a function using the `corr.test` results to highlight the strongest relationships.

Let's build our pipeline function, discard all results except for the significance (p-value) and correlations, and set our cut-off thresholds:

```

Get_Top_Relationships <- function(data_set, correlation_abs_threshold=0.8,
                                pvalue_threshold=0.01) {

  require(psych)
  require(dplyr)

  feature_names <- names(data_set)

  # strip var names to index for pair-wise identification
  names(data_set) <- seq(1:ncol(data_set))
  # calculate correlation and significance numbers
  cor_data_df <- corr.test(data_set)

  # apply var names to correlation matrix over index
  rownames(cor_data_df$r) <- feature_names
  colnames(cor_data_df$r) <- feature_names

  # top cor and sig
  relationships_set <- cor_data_df$ci[,c('r','p')]

  # apply var names to data over index pairs
  relationships_set$feature_1 <- feature_names[as.numeric(sapply(strsplit(rownames(relationships_set), "-"), `
[, 1]))]
  relationships_set$feature_2 <- feature_names[as.numeric(sapply(strsplit(rownames(relationships_set), "-"), `
[, 2]))]

  relationships_set <- select(relationships_set, feature_1, feature_2, r, p) %>%
    rename(correlaton=r, pvalue=p)

  # return only the most insteresting relationships
  return(filter(relationships_set, abs(correlaton) > correlation_abs_threshold |
    pvalue < pvalue_threshold) %>% arrange(pvalue))

}

dim(Get_Top_Relationships(mtcars))

```

```
## [1] 39 4
```

```
head(Get_Top_Relationships(mtcars))
```

```
##   feature_1 feature_2 correlaton      pvalue
## 1      cyl      disp  0.9020329 1.803002e-12
## 2     disp        wt  0.8879799 1.222311e-11
## 3      mpg        wt -0.8676594 1.293958e-10
## 4      mpg      cyl -0.8521620 6.112688e-10
## 5      mpg     disp -0.8475514 9.380328e-10
## 6      cyl        hp  0.8324475 3.477861e-09
```